

The Prompting Handbook

The 2026 guide to the frontier image and video models: what each one is best at, the ideas behind how they generate and edit, and how to prompt them for product, marketing and cinematic work.

Contents

01 Prompting for image	03
Model comparison	04
Image generation	05
Text in the image	09
Showcase: Recraft V4.1	11
Image editing	12
Field notes: killing the slop · controlling the frame	15
Field example: character sheets	17
Field notes: words that work	18
02 Prompting for video	19
Model comparison · features & key concepts	20
Start/end frame vs reference images	22
Video generation	23
Video editing	25
Pipeline: rewrite the first frame	27
Coverage: one location, every angle	28
03 Directing the Shot	29
Field notes: engine rules · the five seams	30
Field notes: scene rules	31
Twelve camera moves	32
Camera & lens language · shot sizes & angles	33
Dialogue · beat-by-beat timing	34
The audition · performance references	35
APPENDIX Cheat sheet	36

COMPANION SPACE · A Space is a shared Magnific canvas where prompts and their results sit side by side, ready to re-run. Every example in this handbook lives on one. [Open the Space](#).



PART 01

Prompting for image

The scene in plain language; the model builds it.

Today's image models work differently from the old keyword piles: describe the scene in plain language and the model builds it. The models aren't interchangeable; each is built for a different job, so choosing the right one matters as much as the words. Here are the seven strongest in Magnific and what each is best at.

Model comparison

Model	Best for	References	Max res
Nano Banana Pro GOOGLE · GEMINI 3 PRO	Cinematic & aesthetic-led work, the hardest prompts, hero imagery	Style, character, product, image	4K
Nano Banana 2 GOOGLE · GEMINI 3.1 FLASH	Editing, product & brand work, high-volume iteration	Style, character, product, image	4K
Seedream 5.0 BYTEDANCE	Posters, packaging, e-commerce, brand-consistent batches	Style, character, product, image (up to 10)	4K
GPT Image 2 OPENAI	Layout-heavy design, infographics, text-accurate visuals	Style, character, product, image	4K
Recraft V4.1 RERAFT	Logos, icons, vector assets, on-brand design systems	Style	SVG · 2K
Flux.2 BLACK FOREST LABS	Brand-exact colour, structured shot-list briefs	Style, character, product, image (up to 10)	2K
Krea 2 KREA	Aesthetic-first looks from style references and moodboards	Style, image	2K

Choosing between models

Nano Banana 2 (Gemini 3.1 Flash Image, Feb 2026) is the newer model and handles most everyday jobs. It offers near-Pro quality at higher speed, reliable text rendering, and the strongest multi-reference consistency. It is the default for editing, product and brand work.

Nano Banana Pro (Gemini 3 Pro Image, Nov 2025) is the premium pick when you need accuracy, brand precision, or when refining the most difficult prompts. It takes longer time to process but typically reaches a higher end result in fidelity. The two cost the same per image, so route by behaviour, not price.

GPT Image 2 is best for complex layout-focused designs.

Recraft earns its slot for typographic dependent posters and art-school level illustration, and is the only current option that exports vector files.

Flux.2 is built for control. It accepts up to ten references at once, matches exact hex colours, and follows structured, shot-list prompts, which suits locked-down brand and product pipelines. Give it an explicit brief; a vague prompt returns a polished but generic image.

Krea 2 is built for style. It works from style references and moodboards and produces looks that avoid the default AI polish. The trade-off is precision: Krea 2 interprets the brief rather than following it exactly, and it is weak at in-image text.

Image generation

How you write the prompt now matters more than any quality tags you add: describe a coherent scene and the model follows it. The reasoning models (**Nano Banana Pro** and **2**, **GPT Image 2**, and the **Seedream** line) read a busy prompt in parts before rendering, so they hold up when a scene has many subjects or a lot of detail.

COMPANION SPACE · [Open the image creation canvas](#) to run every example in this section live.

PROMPT STRUCTURE · IMAGE

Part	What it answers	In the prompt
Subject	Who or what is in frame	a matte-black cologne bottle
Action / state	What it is doing, or its condition	a single droplet hits the slate in a crown splash
Setting	Where, and on what	on wet slate, charcoal background
Lighting	Name the type of light	hard raking light from the right, deep shadows
Composition	Where things sit in the frame	right third of frame, tight macro framing
Style	The finish	photorealistic

- **Scene, not tags.** State the subject and what it's doing, where, in what light, framed how, in which style. Skip boosters like “8K” or “masterpiece.”
- **Quote in-image text.** Put exact words in quotation marks so they render correctly; specify placement (“title top-centre”). The same goes for product copy: repeat the label or pack text verbatim in the prompt, or the model will invent its own.
- **Iterate cheap, deliver sharp.** Explore at 1K, then re-run the keeper at 2K–4K.
- **Lock identity with references** when you need the same character or product across a set.
- **Real products are the hard test.** Details, textures and label type have to survive the render. **Nano Banana 2** holds up best for images (generating and editing alike) and **Seedance 2.5** for video. Zoom the label before you ship.

Product shot

The client wants the fragrance to feel alive, not posed: energy on the shelf. The shot freezes a droplet mid-splash beside the bottle in dramatic light. The prompt names the moment, the light and the framing, and nothing else.



PROMPT

A matte-black cologne bottle on wet slate, right third of frame; a single droplet hits the slate beside it in a small crown splash, fine mist, hard raking light from the right, deep shadows stretching left, charcoal background, tight macro framing, photorealistic.

Layout

Marketing wants an ad that reads from across the room, so the message is the design and the product plays a supporting role. The prompt is a layout brief: what the text says, where each element sits, what stays empty.

MORNINGS, UPGRADED

**Better coffee.
Built for your daily grind.**



PROMPT

A bold vertical social ad, flat editorial print design. The headline “MORNINGS, UPGRADED” enormous in a heavy ink-black grotesque, filling the top half across two lines. Small sub-line “Better coffee. Built for your daily grind.” A brushed steel coffee tumbler bottom-right with a soft shadow, one thin cobalt-blue rule as the only accent, clean white background, generous negative space, no other text.

The long brief



PROMPT

Editorial portrait photo with a vintage 1970s vibe. Young man with a thin mustache and small goatee, short hair, nose piercing, wearing oversized aviator sunglasses with amber lenses, a brown checkered suit jacket and a burgundy shirt. Extreme close-up with a dramatic top-down angle, wide-angle/fisheye look (14-18mm) with slight facial and hand distortion. He holds a large turquoise ceramic mug very close to the camera; visible gold rings and small finger tattoos. Setting: small retro bathroom with teal/sea-green tile walls, metal lockers or panels on one side, a white door behind, and a high window letting in natural light. Lighting: soft natural window light, gentle shadows, warm tones, moderate contrast. 35mm film look, subtle analog grain, warm-teal color grading. Sharp focus on face and mug, background slightly blurred. Intimate composition, confident attitude, street-fashion/editorial styling, high resolution, ultra realistic.



The idea: a frame that feels like a still from a film, not an AI render: a woman in red crossing a wet street at dusk, real motion blur, muted colour. The look words (Kodak palette, natural grain) do the heavy lifting.

PROMPT

35mm film still: a woman in a red coat dashes across a rain-wet crossing at dusk, mid-stride with motion blur, a bus streaking past behind, sodium streetlight and shop-window glow, muted Kodak palette, natural grain, subject on the left third, full frame edge to edge, no letterboxing, no black bars.

Text in the image


Words are where most prompts fall apart, and where new users flub it constantly. The trick is that models place text the way a designer places it, not the way a writer types it: brief the layout, quote the words.

What works

Put every word in quotation marks, exactly as it should read, and say where it goes (“headline across the top”). Brief zones, columns and hierarchy like a designer. Keep it to a few text elements at one or two sizes. When a word keeps mangling, spell it out letter by letter, e.g. “the label reads STUDIO: S, T, U, D, I, O.” **Seedream** and **GPT Image 2** are the strongest; **Recraft** holds the longest text reliably.

What fails, and why

Unquoted text gets paraphrased: the model treats it as a theme, not a script. Long paragraphs and many small labels degrade into shapes, because text is drawn, not typeset. Don't expect an exact font or kerning; fix a stubborn word in an edit pass instead of re-rolling the whole design.



DESIGN · LAYERED TYPE · NANO BANANA PRO

The idea: the brief asks for a poster where type and botany share one space with real depth. The word is quoted and the layers are directed: some flowers sit in front of the letters, others behind. Quote the exact text, then describe what overlaps it.

PROMPT

Editorial botanical poster with the exact word “NATURE” in enormous bright green uppercase, an elegant retro serif display face, filling the upper two thirds. Deep black background with a dark olive undertone and a dense shadowy botanical layer behind the letters: dark leaves, twisting stems, faded vines. In front, vintage hand-painted flowers (ivory lilies, muted orange marigolds, soft pink roses) overlap parts of the word while other blooms stay behind the letters, a layered collage with real depth. Aged botanical-encyclopedia illustration style, lifted shadows. Poetic, mysterious, organic meets futuristic type.



DESIGN · IN-IMAGE TEXT · NANO BANANA 2

The idea: the client wants the quote itself to carry the poster: the whole aphorism set giant in two blocks filling the frame, the laughing figure overprinted in black, and the whole look named in print terms (screenprint, halftone, risograph, 2-color). Quote the words and the model sets them.

PROMPT

Vintage screenprint design poster, giant green condensed text filling the canvas: "A USER INTERFACE IS LIKE A JOKE." across the top half and "IF YOU HAVE TO EXPLAIN IT, IT'S NOT THAT GOOD." across the bottom half, simple black doodle figure centered in overprint, laughing with head thrown back, one hand slapping a knee, the other holding a computer mouse, halftone dot texture, risograph print effect, distressed ink, off-white paper background, bold graphic design, 2-color print, retro illustration.

Where this leads. When the image is close but not right, stop re-rolling and start editing: name the one thing that should change and pin everything else. That workflow is the next section.

Recraft V4.1

The poster and typography specialist. Layered display type, trained brand styles, and the only route to true vectors.

Full prompt: magnific.com/app/creation/8vVH6JKIrU

Creator: magnific.com/app/creator/pmartin



Take it out

Image editing

Editing is now conversational: you type the change instead of masking by hand, and each model layers on its own specialism.

COMPANION SPACE · [Open the image editing canvas](#) to run every example in this section live.

IN SHORT · Say what to change and what to keep in one plain sentence, e.g. “replace the background, keep the subject exactly.” Name only what should move.

Edit one region, or relight the scene

Nano Banana 2 is the everyday editor: it selects and refines any region, shifts camera angle, focus or colour grade, and relights a scene from day to night. Reach for **Pro** only when the edit needs its increased refinement.

Combine several references into one

Seedream merges up to ten references while preserving faces, lighting and tone; **Nano Banana Pro** keeps up to five characters and fourteen objects coherent.

Get true vectors and on-brand art

Recraft outputs true editable vectors (SVG), not just PNG, and enforces a trained brand style. **Magnific** can also trace any raster image into SVG, plus upscale, in/outpaint and remove backgrounds.

Make instructions land exactly

Put exact wording in quotes for text edits, and name only what should change so the rest of the frame stays put. Weak: “fix the headline.” Strong: “Change the headline text only, to ‘FOCUS’. Keep typeface, size, colour and position.”

EDIT · BACKGROUND SWAP · NANO BANANA 2 · BEFORE → AFTER



The idea: take the product exactly as shot and move it into a studio: same bottle, same shadow, new world behind it. One sentence states the change; one states what must not move.

PROMPT

@image1 Keep the product and its shadow exactly; replace the background with a soft, sunlit grey studio sweep.

COMPOSE · MULTI-REFERENCE · NANO BANANA 2



The idea: put this person together with this product in a scene that never happened: the model (@image1) and the tumbler packshot (@image2) composed into one sunlit moment, face and label preserved. Label your references and address them directly.

PROMPT

@image1 @image2 The model from @image1 leans against a sunlit concrete wall raising the tumbler from @image2 toward camera, laughing off-camera, golden side light, shallow focus. Keep her face and the tumbler's label exactly.

Edit · two directions from one original · Nano Banana Pro

The idea: one original, two independent edits: name only the change and pin everything else. Both prompts ran against the same start image; nothing else moved.



PROMPT · Change the time of day to midday. Keep everything else intact.



PROMPT · Add a person sitting up against the tree. Keep everything else intact.

Field notes: killing the slop

Left alone, every model drifts toward the same default: subject centred and staring at the lens, posed and evenly lit, everything in focus, with a processed gloss creeping in on top. There's a word for it now, and your audience knows it: slop. People clock it in half a second, and it tells them a brand didn't care enough to look. Never ship it. Slop is the default, so every parameter you set (the framing, the moment, the lens, the time of day, the light) moves the image one step away from it; every choice you leave open, the model fills with its most average answer. The fix is two habits: direct the shot (off-centre framing, a real moment, depth) and then finish the grade toward soft cinematic glass: gentle contrast, creamy focus roll-off, fine uniform grain. The two blocks below handle the grade; the framing is on you to name.

THE FILM-LOOK LINE

soft cinematic lens, shallow depth of field with creamy focus roll-off, gentle contrast with soft highlight roll-off, fine uniform film grain, no oversharpener, no digital look

THE SKIN-REALISM BLOCK

Natural skin texture: realistic pores, fine hairs, subtle veining, a faint highlight. Tactile and real, not retouched plastic.

Banned words. These summon the slop. Don't write them: ultra sharp · hyper detailed · crisp · razor sharp · 8K clarity · HDR · ultra-realistic detail. Grain is the one film artifact worth keeping; never dust, scratches or “old damaged film.”

UNDIRECTED → DIRECTED · SAME MODEL, SAME SUBJECT · NANO BANANA 2



The idea: the left is the real default: the one-line prompt “A smiling woman holding a basket of vegetables at a farmers market stall.” and nothing else. Clean and competent, but undirected: subject centred, facing the lens, posed smile, everything in focus. Instantly readable as AI stock. The right is the same woman with the shot directed: right third, caught laughing off-frame, blurred foreground framing the shot, warm side light, film grain. Count the choices below: each one is a step away from the default.

PROMPT · DIRECTED

run on the left image as reference · Re-compose the reference as a directed documentary photograph. She is off-centre on the right third, turned three-quarters away, caught mid-movement arranging produce and laughing at someone off-frame. Blurred crate and vegetables in the near foreground on the left; shallow depth of field, the market melts into bokeh. Late-afternoon side light from the left, warm, soft rim light on her hat. Natural matte skin with visible pores and flyaway hairs. 35mm reportage film look, muted palette, faded blacks, gentle contrast, fine grain, no HDR.

Field notes: controlling the frame

The model fills every gap you leave open. These four habits close the gaps that pull a frame off-brief most often.

Word leakage

A mood word colours the whole frame, not just the thing it describes. Write “dewy” or “sensual” anywhere, even about the lighting, and the model nudges the entire image toward suggestive: clothing starts slipping and expressions turn pouty. Counter it with explicit instructions, even when they feel unnecessary. Straps stay neatly in place; lips are closed and relaxed.

Negative prompts

Keep a standing ban list for the usual drift and paste it into every portrait prompt: smiling, glamour makeup, jewelry, plastic smooth skin, dramatic rim light, colored background, soft focus. Escalate only when needed: no pouting, no parted lips.

Spell out stubborn words

When rendered text keeps mangling a word (“FOIUS” for FOCUS), restate it letter by letter in the edit prompt: “spelled F, O, C, U, S”, and it lands. The “SIGNAL” program poster earlier in this part was fixed exactly this way.

Character sheets

For recurring characters, one three-panel sheet becomes the master reference every later generation is built from. It earns its own section, on the next page.

Character sheets

For recurring characters nothing pays off more: build a three-panel character sheet before a character appears twice anywhere, and make it the master reference. Every later generation takes the sheet along as a reference; every prompt describing the character is written from it. When the identity already exists, start from a reference portrait, as here; when it doesn't, generate the sheet from text.

FIELD EXAMPLE · CHARACTER SHEET · NANO BANANA 2 · FROM A REFERENCE



The idea: turn one portrait into the character's official ID: straight-on, side profile, full body in the final wardrobe, all on one sheet. From here on, the sheet rides along as a reference in every generation, and every prompt describing her is written from it.

PROMPT

reference: her portrait · Character sheet: one 16:9 frame divided into three equal panels on a pure white studio background. Left panel: straight-on head-and-shoulders portrait. Centre panel: exact side profile portrait. Right panel: full body, standing relaxed. The same woman in all three panels, the woman from the reference image, keeping her identity exactly: dark hair slicked back, magenta acetate sunglasses, natural lips, blue technical shell jacket with orange trim details, matching blue trousers and clean sneakers. Soft even studio light, no props, no text, fine film grain, photorealistic.

When the character changes, the sheet changes first. New wardrobe, new haircut, new age: edit the sheet, not the shots. Regenerate it once with the change described, then keep prompting from the new sheet; a character has exactly one current sheet at any time.

Field notes: words that work

Half of prompting is knowing the name of the look. Teams often can't get the image they want simply because they don't have the vocabulary for it, so here is the tested shortlist. These are image-generation keywords; stack two or three per prompt, never all of them: stacked look words dilute each other, and the strongest one stops landing.

Light

The most reliable lighting keyword in testing is shadowplay: natural, sophisticated light with real shadow structure. Pairs beautifully with cinematic.

shadowplay · single soft key from upper left · hard shaft of light · practical light only · low-key · rim light · golden hour

Grade & mood

Two or three mood words maximum; past that they cancel each other out.

muted palette · faded blacks · gentle contrast · cinematic · editorial · quiet-luxury · understated

Texture

Texture words are what separate a photograph from a render.

fine uniform film grain · visible pores · tactile fabric weave · atmospheric haze · soft highlight bloom

Composition

Name where the empty space goes and the layout follows.

generous negative space · centered, deadpan framing · chest-up crop · slight three-quarter angle · leading lines

Blur & focus: say the effect, not the gear

These do what people hope f-stops will do.

background melts into soft blur · creamy focus roll-off · subject razor sharp against a soft distance · only the label in focus

Stacks that work

Tested pairings, in order: sophisticated natural light, quiet-luxury product, lived-in interiors.

shadowplay + cinematic + fine film grain · muted palette + faded blacks + negative space · practical light only + atmospheric haze

Gear numbers are weak levers. In team testing across nearly every model, focal lengths and apertures (35mm, 50mm, f/1.2) barely register: models don't differentiate these technical settings yet (GPT Image 2 is the partial exception). That isn't a contradiction of the "35mm film still" example earlier: there the lens number is paired with visual words that do the real work: motion blur, muted Kodak palette, natural grain. Keep gear numbers if you like, but always pair them with the visible effect you want, because the effect language is what the render actually responds to.

A painterly illustration of a superhero character, likely Robin, crouching on a rocky ledge. The character has dark hair, wears a black mask, a red cape, a green long-sleeved shirt, and green pants. They are wearing white gloves and brown shoes. The background shows a lush green valley with a small town and a river, viewed from a high vantage point. The style is highly detailed and artistic, with visible brushstrokes.

PART 02

Prompting for video

Motion, sound and performance, generated in a single pass.

DEEP
ART

The 2026 leap is the shift to unified, “omni” architectures: text, image, video and audio go into one model, and finished footage with synchronized sound comes out in a single pass. The frontier is led by Seedance 2.5, with Kling 3.0 Omni, Alibaba's Happy Horse and Google's brand-new Gemini Omni Flash close behind.

Model comparison

Model	Inputs (multimodal)	Native audio	Max res	Best for
Seedance 2.5 BYTEDANCE	Text, image, video, audio (up to 50 mixed image & video references, plus audio), plus character, product, style	Yes · lip-sync	30s · 4K	The default; best quality and aesthetics, references, audio, narrative
Kling 3.0 Omni KUAISHOU	Text, image, video (performance cloning of face & voice)	Yes · lip-sync	15s · 4K	Character continuity, action, dialogue, accurate on-product text
Happy Horse 1.1 ALIBABA	Text, image (up to 9); video-to-video & subject-swap editing	Yes · lip-sync	15s · 1080p	Cinematic short-form, portrait & talking-head, fast iteration
Gemini Omni Flash GOOGLE · NEW	Text, image, video (audio input not yet)	Yes · own voice only	10s · 720p	Conversational direction; matches Seedance on adherence, trails on quality

Choosing between models

Seedance 2.5 is the default for anything that has to ship. It has the best quality and aesthetics, native audio, and clips up to 30 seconds. Its reference system accepts up to fifty images and videos in one prompt.

Gemini Omni Flash matches it on prompt adherence but trails clearly on quality and look (it also generates its own sound and speech, but won't sync lips to an audio file you supply): direct loosely there, finish in Seedance.

Kling 3.0 Omni earns its slot when performance cloning, physics or on-product text carry the shot; **Happy Horse** when credits are tight.

Sora 2 and the Wan line round out the wider field.

Features & key concepts

Six capabilities define the omni generation of video engines, and the rest of this part builds on them. Two of the six are confused so often that they get their own comparison below.

Multi-modal (“omni”)

One model ingests text, images, video and audio together and outputs video with sound in a single pass. This single-pass design is the biggest change from the previous generation of video engines, which took a text prompt and returned silent picture.

Native audio

Dialogue, effects and ambient sound are generated alongside the picture and locked to it, rather than added in post, including lip-synced speech.

Start / end frame

Supply a first frame (and optionally a last) and the model animates the motion between them, giving precise control over how a shot opens and lands.

References

Image references fix a character or product; video references teach motion and voice; audio references set a voice or soundscape.

Multi-shot storyboarding

Several cuts are generated in one job with continuity, each shot carrying its own prompt for duration, framing and action.

Character consistency

“Elements” and coreference lock identities so multiple characters stay distinct and on-model across every shot and scene change.


Start/end frame vs reference images

These two controls are easy to conflate, but they answer different questions, and you can use both in the same generation.



Start / end frame: controls the moment

You pin the first frame (and optionally the last) of a single shot, and the model generates the motion between them. They govern how this one clip opens and lands. A different shot means different frames.



Reference images: control the identity

You show the model what a character, product or style looks like, so it stays consistent at any point in the timeline and across many shots. They don't fix a moment; they inform appearance.

IN SHORT · Start/end frames decide where the shot begins and ends; references decide who or what appears, and in what style. Pin a start frame to control the opening composition, and attach a character reference so that character stays on-model, in the same generation.

Video generation

All four support text-to-video and image-to-video; reference-to-video and frame-to-video vary by model (see the table). The prompt is a shot description, not a bag of tags: think like a director handing instructions to a camera crew.

COMPANION SPACE · [Open the video generation canvas](#) to run every example in this section live.

PROMPT STRUCTURE · VIDEO

Part	What it answers	In the prompt
Shot	The framing	vertical 9:16 medium shot
Camera move	One move per clip	quick push-in
Subject	Who or what	a hand and a coffee flask
Action	One action per clip	lifts the flask off the counter
Setting	Where, and when	a sunlit kitchen counter, morning
Lighting	Name the type of light	bright natural morning light
Style	The finish	clean and airy
Audio	What is heard	SFX: gentle pour; no music, no subtitles

Each model orders it slightly differently. **Seedance:** state shots, duration and aspect ratio up front, then subject → action → environment → camera → style, labelling references as @image1 / @video1 / @audio1. **Kling:** subject + movement + scene + camera + lighting + atmosphere. **Happy Horse** follows the same shot-description discipline. **Gemini Omni Flash** is the exception: it is tuned for conversational direction, so you can prompt loosely and refine over follow-up turns.

- **Front-load the shot.** Open with the shot type and one camera move; the model reads the first line first.
- **One action, one camera move per clip.** Stacking either produces warping and instability.
- **Use real cinematography language.** “Slow dolly-in,” named light sources, exact shot sizes. Concrete moves and light beat mood words, and lens numbers barely register.
- **Control the audio.** SFX and ambient in plain words, dialogue in quotes. Write **“No music”** to get clean SFX only, then add your track when editing; add “no subtitles” to avoid burned-in captions.
- **Repeat product copy verbatim.** Any on-pack or on-screen text must be quoted in the prompt, exactly as it should read, in video as in image. **Kling 3.0** is the strongest at keeping product copy correct in motion.
- **For image-to-video, describe motion and camera only;** the frame already supplies the subject.
- **Iterate cheap, finish sharp.** Lock composition at 720p/1080p, then re-render the keeper at 4K.



PRODUCT SHOT · IMAGE-TO-VIDEO · SEEDANCE 2.5

The idea: bring a product still to life as a hero shot: one slow orbit around the sneaker, a rim-light sweep across the sole. One move, one clip; SFX only so a track can be scored later.

PROMPT · MOTION + CAMERA ONLY

@image1 Slow 360° orbit around the sneaker from @image1, a rim-light sweep across the sole, soft studio cyclorama, shallow focus. 5 seconds. SFX: soft room tone. No music. No subtitles.



MARKETING · SOCIAL AD · SEEDANCE 2.5

The idea: a vertical beat for social in the handheld, user-generated (UGC) style: a hand lifts the flask off a sunlit counter, quick push-in, rising steam. Bright and airy, captions kept clear, no music so a track drops in cleanly.

PROMPT

Vertical 9:16 medium shot. A hand lifts a coffee flask off a sunlit kitchen counter, quick push-in, steam rising. Bright natural morning light, clean and airy. SFX: gentle pour, ambient kitchen. No music. No subtitles.



CINEMATIC · TEXT-TO-VIDEO · SEEDANCE 2.5

The idea: a cinematic mood piece: a low side-tracking shot travelling with a cyclist on a coastal road at golden hour, grass bending in the wind. SFX only (tyres, gulls, wind) so it can be scored later.

PROMPT

Medium tracking shot. A cyclist coasts down a coastal road at golden hour, camera moving alongside at low height, grass bending in the wind. Shallow focus, warm film grade, fine grain. SFX: tyres on tarmac, distant gulls, soft wind. No music. No subtitles.

Why every prompt here ends “No music. No subtitles.” Generated music locks you in and burned-in captions can't be removed, so generate clean: SFX and ambience only, captions off. Score and type are added in the edit, where you control them. If a clip carries dialogue, keep “no subtitles” and let the spoken line live in quotes instead.

Video editing

Editing is where the generation-versus-refinement split shows. Generation models (Seedance 2.5, Happy Horse) are strongest when you create a new clip and visual impact matters most. Editing models (Kling 3.0 Omni, Gemini Omni Flash) are strongest when you change an existing clip and staying faithful to the source footage matters most. Score the music last; leave it out at generation time. Below are the most common operations:

Video-to-video & restyle	Re-skin or remaster existing footage. Kling 3.0 Omni's editing pipeline runs 3–15s at 4K in and out, staying faithful to the source.
Extension	Chain generations (each new clip continues from the last frame of the previous one) to push past the base clip length; Seedance and Kling extend a shot in place while holding visual continuity.
Reframe & upscale	Change aspect ratio for a new platform, and upscale for delivery: the generation models top out at 4K natively, and dedicated upscale tools reach up to 8K.
Voice & audio edits	Swap or dub the spoken voice while preserving timing and visuals, or add the music you held back at generation time.

IN THE FIELD · GEMINI OMNI FLASH

The model re-derives the physics

Ask it to turn a tabletop into a shallow pool and it reasons through the whole scene: the hand reads as wet, water ripples and refracts, shadows and reflections update, and the sound effects follow. That physical consistency is what makes it strong inside edit pipelines.

Change the table to be a shallow pool of water.

IN THE FIELD · GEMINI OMNI FLASH

The model listens to the footage

The model can act on speech inside the clip: the video's own audio becomes part of the instruction. Point it at a line of dialogue and let that drive the edit, no manual transcription needed.

Change the background to whatever colour the subject in the video says.

EDIT · EXTEND · SEEDANCE 2.5




The idea: the shot is good but too short, so make it continue: seven more seconds, seen low from behind, coat swinging, striding toward the arched doorway. No visible cut.

PROMPT

@video1 Extend @video1 by 7 seconds; keep the character, wardrobe and lighting; continue the walk toward the doorway. No cuts.

EDIT · REFRAME & RESTYLE · KLING 3.0 OMNI



The idea: the same shot, re-composed for a vertical feed: the skater held in the lower third with sky above, warm filmic grade, original motion preserved.

PROMPT

@video1 Reframe @video1 to 9:16 for Reels, keep the skater in the lower third with sky above; grade to a warm filmic look; keep the motion.

Pipeline: rewrite the first frame

The strongest video edits often aren't made in a video model at all. Because a clip is driven by its start frame, you can pull frame one out of any video, edit it as an image with **Nano Banana Pro** (where control is surgical) and hand it back to **Seedance 2.5** with the original clip as a motion reference. The result: the same shot, same move, same pacing, with the change you directed. This exact pipeline produced the frames below.



Start with a clip

A 5-second **Seedance 2.5** shot: slow 180° orbit around a white sneaker on a pedestal, rim-light sweep, one continuous take. Any clip works, generated or filmed.



Extract frame one

Export the first frame as a still. In Magnific the start frame of every generation is already saved alongside the clip; no tooling needed, grab it and treat it as an image.



Edit the frame with Nano Banana Pro

Make the change as an image edit, where instruction-following is strongest. Name what changes; pin everything else.

PROMPT · NBP

@image1 Change the white sneaker to a matte-black colorway with black laces and a black sole; keep the pedestal, the lighting, the camera angle and everything else exactly.



Re-render with Seedance 2.5

Set the edited frame as the start frame and attach the original clip as a motion reference: the shot repeats its own move around the new subject.

PROMPT · SEEDANCE 2.5 · START FRAME + VIDEO REFERENCE

@image1 @video1 Repeat the exact same slow 180° orbit and rim-light sweep as @video1, now around the matte-black sneaker from the start frame. Same pacing, same framing, one continuous take. No cuts.

Why it works. The start frame pins what the shot opens on; the video reference pins how it moves. Splitting the edit this way gives you image-grade control over content and video-grade continuity over motion, and it chains: edit any frame, extend from it, or restyle a whole sequence one start frame at a time.

Coverage: one location, every angle

A single still of a location can become full scene coverage. Drop the image into **Seedance 2.5** (or **Kling** at 4K) as the start frame and prompt a slow fly-through that rotates and holds on each angle. The render works like a location scout: every held beat is a frame you can extract (matching light, matching set), ready to serve as start frames, references or backdrops for the actual shots.

FIELD EXAMPLE · SCENE COVERAGE · NANO BANANA 2 → SEEDANCE 2.5



The idea: the production needs a location scout for a stadium that exists only as one still. The still becomes the start frame; the fly-through does the rest.

PROMPT · VIDEO · SEEDANCE 2.5

start frame: stadium still · Real-estate style drone fly-through shot from the stands of a packed football stadium during an evening match. The camera slowly moves through the tier, rotating and stopping for a beat facing each angle: the camera rotates right and slowly pushes in over the crowd toward the floodlit pitch, then rotates left and slowly pushes in along the rows of cheering fans close to the lens, then rotates 180 degrees and slowly pushes in toward the upper tier and roofline behind.

Why it works. Name each rotation and give it a stop (“rotates right and slowly pushes in”) rather than asking for “multiple angles.” The held beats are what make frames extractable. And because every angle comes from one continuous render, the crowd, the light and the architecture stay consistent across all of them, which is exactly what generating each angle as a separate image can't guarantee.

PART 03

Directing the Shot

Camera, lens and performance: the language that turns a description into a directed shot.



Field notes: engine rules

Video engines share a set of behaviours worth designing around. They can sound like superstition; treat them as law. Every rule here earned its place the first time a render broke it.

COMPANION SPACE · [Open the directing canvas](#) to run every example in this part live.

Chaining clips: the five seams

A multi-clip film is a chain of seams, and each seam answers one question: does the next clip continue this frame, continue this motion, pass through something, or break clean?

Seam	What it is
Frozen handoff	The last frame of clip N becomes the first frame of clip N+1: end on a freezable state and restate it exactly. This is the seam the pipeline on the previous page runs.
Action bridge	A motion starts at the end of one clip and completes at the start of the next from a new angle. Name the motion and its direction in both prompts.
Match cut	Two frames share a shape, position or colour while the scene changes. Name the rhyming element in both prompts.
Portal	The camera pushes into something that fills the frame (a pupil, steam, a doorway) and emerges somewhere new. Ration it: one per film, or it reads as a gimmick.
Hard cut	A clean break; montage logic. The default when nothing needs to connect.

The brief decides the seam. “Seamless / one take / it should flow” → frozen handoff · “cut to / meanwhile / vignettes” → hard cuts · “morphs into / becomes” → match cut. When the brief is silent, hard cut is the default; never portal.

Two rules on top. Work from what rendered, not what was planned: before writing the next clip, extract the actual final frame and rewrite the opening against it. And lock the product: repeat its description word-for-word in every clip's prompt; that repetition is what keeps it consistent across separate generations. Keep chained clips to 30 seconds each.

Scene rules

These are the physics of a generated scene: how the engines track people, objects and space inside a single shot. Write with them and clips hold together; write against them and the seams show.

Rule	In practice
Action is intent plus a named move, never biomechanics.	Engines execute intent; they cannot follow joint-by-joint choreography. "Spinning back kick connects," not "forearm rotates 45° to deflect the hook." Name what happens and let the model stage how it happens.
Emotions are physics, not labels.	A label asks the model to guess a face; a physical action tells it exactly what to render. "Jaw clenches, nostrils flare," not "looks angry." The audition page later in this part runs entirely on this rule.
Only what can be seen or heard.	The prompt describes a camera feed, so smells, thoughts and backstory render as nothing. Translate them into visible evidence: not "the air smells of cedar," but "cedar shavings on the workbench, dust drifting in the light."
Exit-frame means gone.	Anything that leaves the frame is gone for the rest of that shot; never exit and re-enter in one continuous take. Off-screen means nonexistent: show a state change on camera before referencing it.
Track three characters, max, across cuts.	Beyond three, identities start to swap: faces migrate between bodies and wardrobe details trade places. Keep extras generic and unnamed; describe in detail only the characters who must stay recognisable.
Compose around mirrors, never into them.	Reflections break scene geography: the engine has to invent the reversed room, and it usually invents it wrong. Shoot past a mirror or keep it at a glancing angle; never make the reflection carry the shot.
Re-establish the scene after every cut.	Restate positions and facing directions, and keep the camera on one side of the action line (the 180° rule) so characters don't swap sides of frame.
Free-form on-screen text is unreliable.	Declare a text-safe zone ("the right third of the final frame falls to clean darkness") and set type in post. For on-product label type, Kling 3.0 holds best.

Twelve camera moves

Cinematographers learn these in film school; most AI creators never name one, so the camera just sits there. Name the move and the shot comes alive.

Move	What it does
Push-in	Moves toward the subject. Builds tension and intimacy; keep it slow.
Pull-back	Retreats to reveal. Isolation, scale, endings: the reveal shot.
Pan	Horizontal rotation, camera fixed. Suspense lives in what you haven't shown yet.
Tilt	The vertical pan. Tilt up on a hero and they read as powerful.
Tracking	Camera travels with the subject. Energy, forward motion, "you're there."
Arc / orbit	Circles the subject. A product on a clean pedestal takes a full turn; keep arcs around people under ~30°.
Crane / jib	Sweeps vertically on a boom. Grandeur and scale: the god-view.
Zoom	Focal length changes, camera stays put. Flatter than a dolly; fast zoom = music-video energy.
Dolly zoom	Camera and lens move opposite ways. Background warps, subject holds: pure dread.
Whip pan / crash zoom	Extreme speed for transitions. Shock, comedy, scroll-stopping.
Handheld	Natural, unstabilised shake. Add "subtle" or the model goes full earthquake.
Static + angles	Low, high, Dutch, bird's-eye, worm's-eye. Low = power; Dutch = unease; bird's-eye = scale.

The mistake everyone makes is stacking moves into one prompt. One move, one clip, always. And add "slow" to almost everything: slow moves hide what the model can't render cleanly; fast moves expose every flaw.

Camera & lens language

Real optical and framing terms give you a precise shared language, and shot sizes and angles are strong prompt controls. Lens and aperture numbers are weaker: models barely differentiate 35mm from 85mm or f/1.2 from f/8, so treat the specs below as shorthand and always pair them with the visible effect (“background melts into soft blur”) you are actually after.

Lenses & optics

Term	In a prompt	Effect
Wide-angle · 14–24mm	“shot on a 24mm wide lens”	Expansive, immersive; mild edge distortion
Standard · 35–50mm	“35mm”	Natural, close to human vision
Telephoto · 85–135mm	“85mm portrait lens”	Compression; flattering; soft background
Macro	“100mm macro”	Extreme detail; a tiny subject fills the frame
Anamorphic	“anamorphic, 2.39:1”	Widescreen; horizontal flares; oval bokeh
Fast aperture	“f/1.8, shallow depth of field”	Subject isolation; creamy bokeh
Deep focus	“f/11, deep focus”	Sharp from foreground to background
Rack focus	“rack focus to her face”	Shifts attention within a single shot

Shot sizes & angles

Establishing / extreme wide: sets place and scale.	Wide: the full subject in its environment.
Medium: waist-up; the conversational default.	Close-up: a face or detail fills the frame.
Extreme close-up: eyes, a texture, a logo.	Over-the-shoulder / POV: framed past a shoulder, or through the character's eyes.
Low / high angle: up for power, down for vulnerability.	Dutch angle: tilted horizon; unease and tension.
Bird's-eye / top-down: directly overhead; scale and pattern.	Worm's-eye: from the ground up; monumental.

Dialogue

Native-audio models (**Seedance**, **Kling**, **Happy Horse**) will speak the lines you write and lip-sync them; **Gemini Omni Flash** speaks too, but won't sync to audio you supply. Put the words in quotation marks and attribute them to a character, keep each line to what fits the clip, and describe the delivery. One speaker per beat keeps the sync clean; add “no subtitles” to avoid burned-in captions.



DIALOGUE · LIP-SYNC · KLING 3.0 OMNI

The idea: one spoken line, delivered warmly to camera: the barista slides the cup across, says her line, lips synced, café ambience behind. Quote the words, name the delivery.

PROMPT

Close-up of a barista sliding a cup across the counter. She looks to camera and says, warmly, “Careful, it's hot.” Soft café murmur, espresso hiss. No music. No subtitles.

Beat-by-beat timing

For a shot with more than one action, choreograph it with timecodes. Number the beats from 00:00, give each a single action, and keep the last beat inside the clip length. This stops the model cramming everything into the first second and lets you sync action to a product reveal or a musical beat.



TIMING · BEAT SHEET · SEEDANCE 2.5

The idea: choreograph time itself: a 6-second watch reveal in three named beats (rest, rise, hero), one action each, so the model paces the shot instead of improvising it.

PROMPT

6-second product shot, one continuous take. 00:00 the watch lies flat on black stone, still. 00:02 it rises and rotates slowly to face camera. 00:04 it settles as a rim light sweeps across the dial. SFX: low hum, soft click. No music. No subtitles.

The audition: screen-test the performance

The character sheet locks the face; the audition locks the performance. Before a character carries a real scene, run a screen test: a short self-tape that shows how they move, speak and react under direction. Generate two or three takes with different voice recipes, pick the performance, and attach the winning clip as a video reference in every scene that follows, the same way the sheet rides along for identity. Deciding the acting once is far cheaper than re-rolling scenes until it lands by accident.

Shoot a self-tape, not a trailer: simple room, medium close-up, eye-level, shallow focus. The frame exists to reveal acting, not spectacle.

A voice recipe, three qualities max: “exhausted, tender, restrained.” Mix emotional forces; one flat label reads dead.

End unresolved: let silence sit before and after the line. The held moment is where the performance lives.

One short line with subtext: playable more than one way, implying someone off-camera. Scene dialogue, not lore or narration.

Direct the micro-acting: one vocal moment (a broken start, a stressed word, a quiet crack) and one small physical behaviour (a gaze shift, jaw tension, a look down and back).



AUDITION · SELF-TAPE · SEEDANCE 2.5

The idea: a casting-room take: one line with subtext, a three-quality voice recipe, the micro-acting named beat by beat. Run it against two other recipes, keep the take that lands, and it becomes the character's performance reference.

PROMPT

Cinematic self-tape audition, one continuous take. Medium close-up, locked-off eye-level camera, shallow depth of field: a woman in her 30s against a plain grey wall, soft window light from the left, listening to an off-camera reader. A beat of silence; she looks down, then back up. She says, quietly: “You knew. The whole time, you knew.” Voice: controlled, hurt, almost warm; it catches on the second “knew.” Jaw tightens, one slow blink, the trace of a smile that is not one. She holds eye contact and says nothing more. Natural skin texture, fine film grain, muted palette. No music. No subtitles.

AUDITION WORKFLOW ADAPTED FROM @_OAK200

The cheat sheet

Everything injectable, in one place. Pin it.

The film-look line

IMAGE + VIDEO

Kill the processed AI default in every single prompt: soft glass, gentle contrast, fine uniform grain.

EXAMPLE PROMPT

soft cinematic lens, shallow depth of field with creamy focus roll-off, gentle contrast with soft highlight roll-off, fine uniform film grain, no oversharpening, no digital look

Standard negative set

IMAGE

Run the standard negatives to strip the usual drift; escalate only when a render misbehaves.

EXAMPLE NEGATIVE PROMPT

smiling, glamour makeup, jewelry, plastic smooth skin, dramatic rim light, colored background, soft focus · escalate: no pouting, no parted lips, no seductive expression

Model routing

IMAGE + VIDEO

Route by job, not by habit. **Stills:** NB Pro for cinematic and hero · NB 2 for edits, product, volume · GPT Image 2 for layout · Recraft for posters and vectors. **Video:** Seedance 2.5 by default · Kling 3.0 for cloning, physics, on-product text · Happy Horse on a budget.

EXAMPLE

hero perfume still → Nano Banana Pro · replace the background → Nano Banana 2 · gig poster with display type → Recraft V4.1 · 8-second product film → Seedance 2.5

Always · image

IMAGE

Quote in-image text and spell stubborn words letter by letter · lock identity with references · name only what should change in an edit · work from what rendered, not what you planned; rewrite forward from it.

EXAMPLE PROMPT

Change the headline to read "FOCUS", spelled F, O, C, U, S, top-centre, same typeface and size; keep everything else exactly unchanged.

The skin-realism block

IMAGE + VIDEO

Whenever skin is visible, spell out that it is tactile and real; otherwise models retouch it to plastic.

EXAMPLE PROMPT

Natural skin texture: realistic pores, fine hairs, subtle veining, a faint highlight. Tactile and real, not retouched plastic.

Banned words

IMAGE + VIDEO

These summon the slop; never write them. Grain is the only film artifact worth keeping; dust, scratches and “old damaged film” never.

NEVER WRITE

ultra sharp · hyper detailed · crisp · razor sharp · 8K clarity · HDR · ultra-realistic detail

Keywords that work

IMAGE

When the look won't come, the vocabulary is usually missing: reach for the tested stacks.

EXAMPLE STACKS

shadowplay + cinematic + fine uniform film grain for sophisticated natural light · muted palette + faded blacks + generous negative space for quiet-luxury product · background melts into soft blur for what f-stops can't do

Always · video

VIDEO

Chain in clips of 30 seconds or less · three tracked characters max · one portal per film · repeat the product description verbatim in every clip · restate positions and facing after every cut.

EXAMPLE

Clip 2 continues the final frame of clip 1: she stands at the counter facing left, camera over her right shoulder; she lifts the flask and turns to the window.

Credit costs change with models and plans. Current pricing: magnific.com/pricing.

Who's behind this guide

This guide was put together by us at Magnific Studios, a global team of 40+ AI creative experts across Europe, the US and Asia. We're creatives who know the ins and outs of Magnific and every image and video model out there, and our job is to help teams figure out everything these tools can actually do for them, whether that's a film studio, a fashion brand, or any team with a bold creative vision ready to shake things up.

We work alongside our customers end to end. Before they commit, we build proofs of concept tailored to their own use cases, so they can see exactly what's possible. Once they're on board, we help them make it stick: hands-on onboarding, consulting sessions (remote or on-site), and dedicated support to bring their projects to life. The same techniques in this guide are the ones we use every day to produce high-impact work for our clients.

Want to see what Magnific can do for your team?

[Schedule a demo](#)